

Image Processing and Computer Vision

1st George Lancaster
dept. of Computer Science
University of Bristol
Bristol, United Kingdom
qv18258@bristol.ac.uk

2nd Ren Jiang
dept. of Computer Science
University of Bristol
Bristol, United Kingdom
mu18336@bristol.ac.uk


Abstract—This report outlines the tasks completed for Image Processing and Computer Vision assignment one. The assignment has been split into four sub-tasks. In task one, we experiment with the Viola-Jones object detector to detect faces.

I. THE VIOLA-JONES OBJECT DETECTOR


We first annotated an image data set containing sixteen images for ground truth. Each image within the data set contains faces, dartboards or a combination of both. In this first task, we use the Viola-Jones object detector to find faces within the images.



(a) darts4.jpg




(b) darts5.jpg



(c) darts13.jpg



(d) darts14.jpg



(e) darts15.jpg

Fig. 1: Five images from the test data set. Green rectangles have been drawn where the Viola-Jones classifier has detected a face.

The true positive rate for images *darts5.jpg* and *darts15.jpg* are 11 and 1 respectively.

Practically, the true positive rate is difficult to assess because we because we need to define our own ruleset for what counts as a detection. For this task, we have defined a true detection as an area that shares a 50 per cent overlap with a ground truth label. This value was chosen to maximise the f1 score when classifying faces from the test data set. The results varying the percentage overlap can be seen in table x.

TABLE I: All values of percentage overlap gave an identical F1 score, when detecting faces from the test data set.

Overlap Threshold (%)	F1 score
0	0.568
60	0.541
65	0.541
68	0.541
69	0.519
70	0.514
80	0.459

Although the true positive rate can be used to indicate a classifiers accuracy, it does not reflect its true performance. It is always possible to get a 100 per cent detection rate on any classification task as we can select all possible areas of an image, regardless if they contain a target or not. The key to a good classifier is to get a high true positive rate, whilst keeping the false positives minimal. We can use the F1 score to measure the relationship between the precision and the recall of the model. The F1 score can therefore be considered to be a more reliable measure of classifier performance.

The F1 score of the Viola-Jones classifier, when detecting faces in the test data set is

II. THE DARTBOARD DETECTOR

A. Subtask a




Fig. 2: True positive rate plotted against false positive rate, when training the cascade classifier on images of dartboards. Each stage of training has been plotted as its own point.

TABLE II: F1 scores for all images, when detecting for dartboards using a cascade classifier trained on dartboards.

Image Name	F1 Score
dart0.jpg	0.5 00
dart1.jpg	nan
dart2.jpg	0.250
dart3.jpg	1.538
dart4.jpg	0.4 00
dart5.jpg	0.154
dart6.jpg	0.333
dart7.jpg	nan
dart8.jpg	0.160
dart9.jpg	0.222
dart10.jpg	0.400
dart11.jpg	0.222
dart12.jpg	0.667
dart13.jpg	0.200
dart14.jpg	0.087
dart15.jpg	0.667
Average f1 score	0.276

III. INTEGRATION WITH SHAPE DETECTORS

IV. FURTHER IMPROVEMENTS